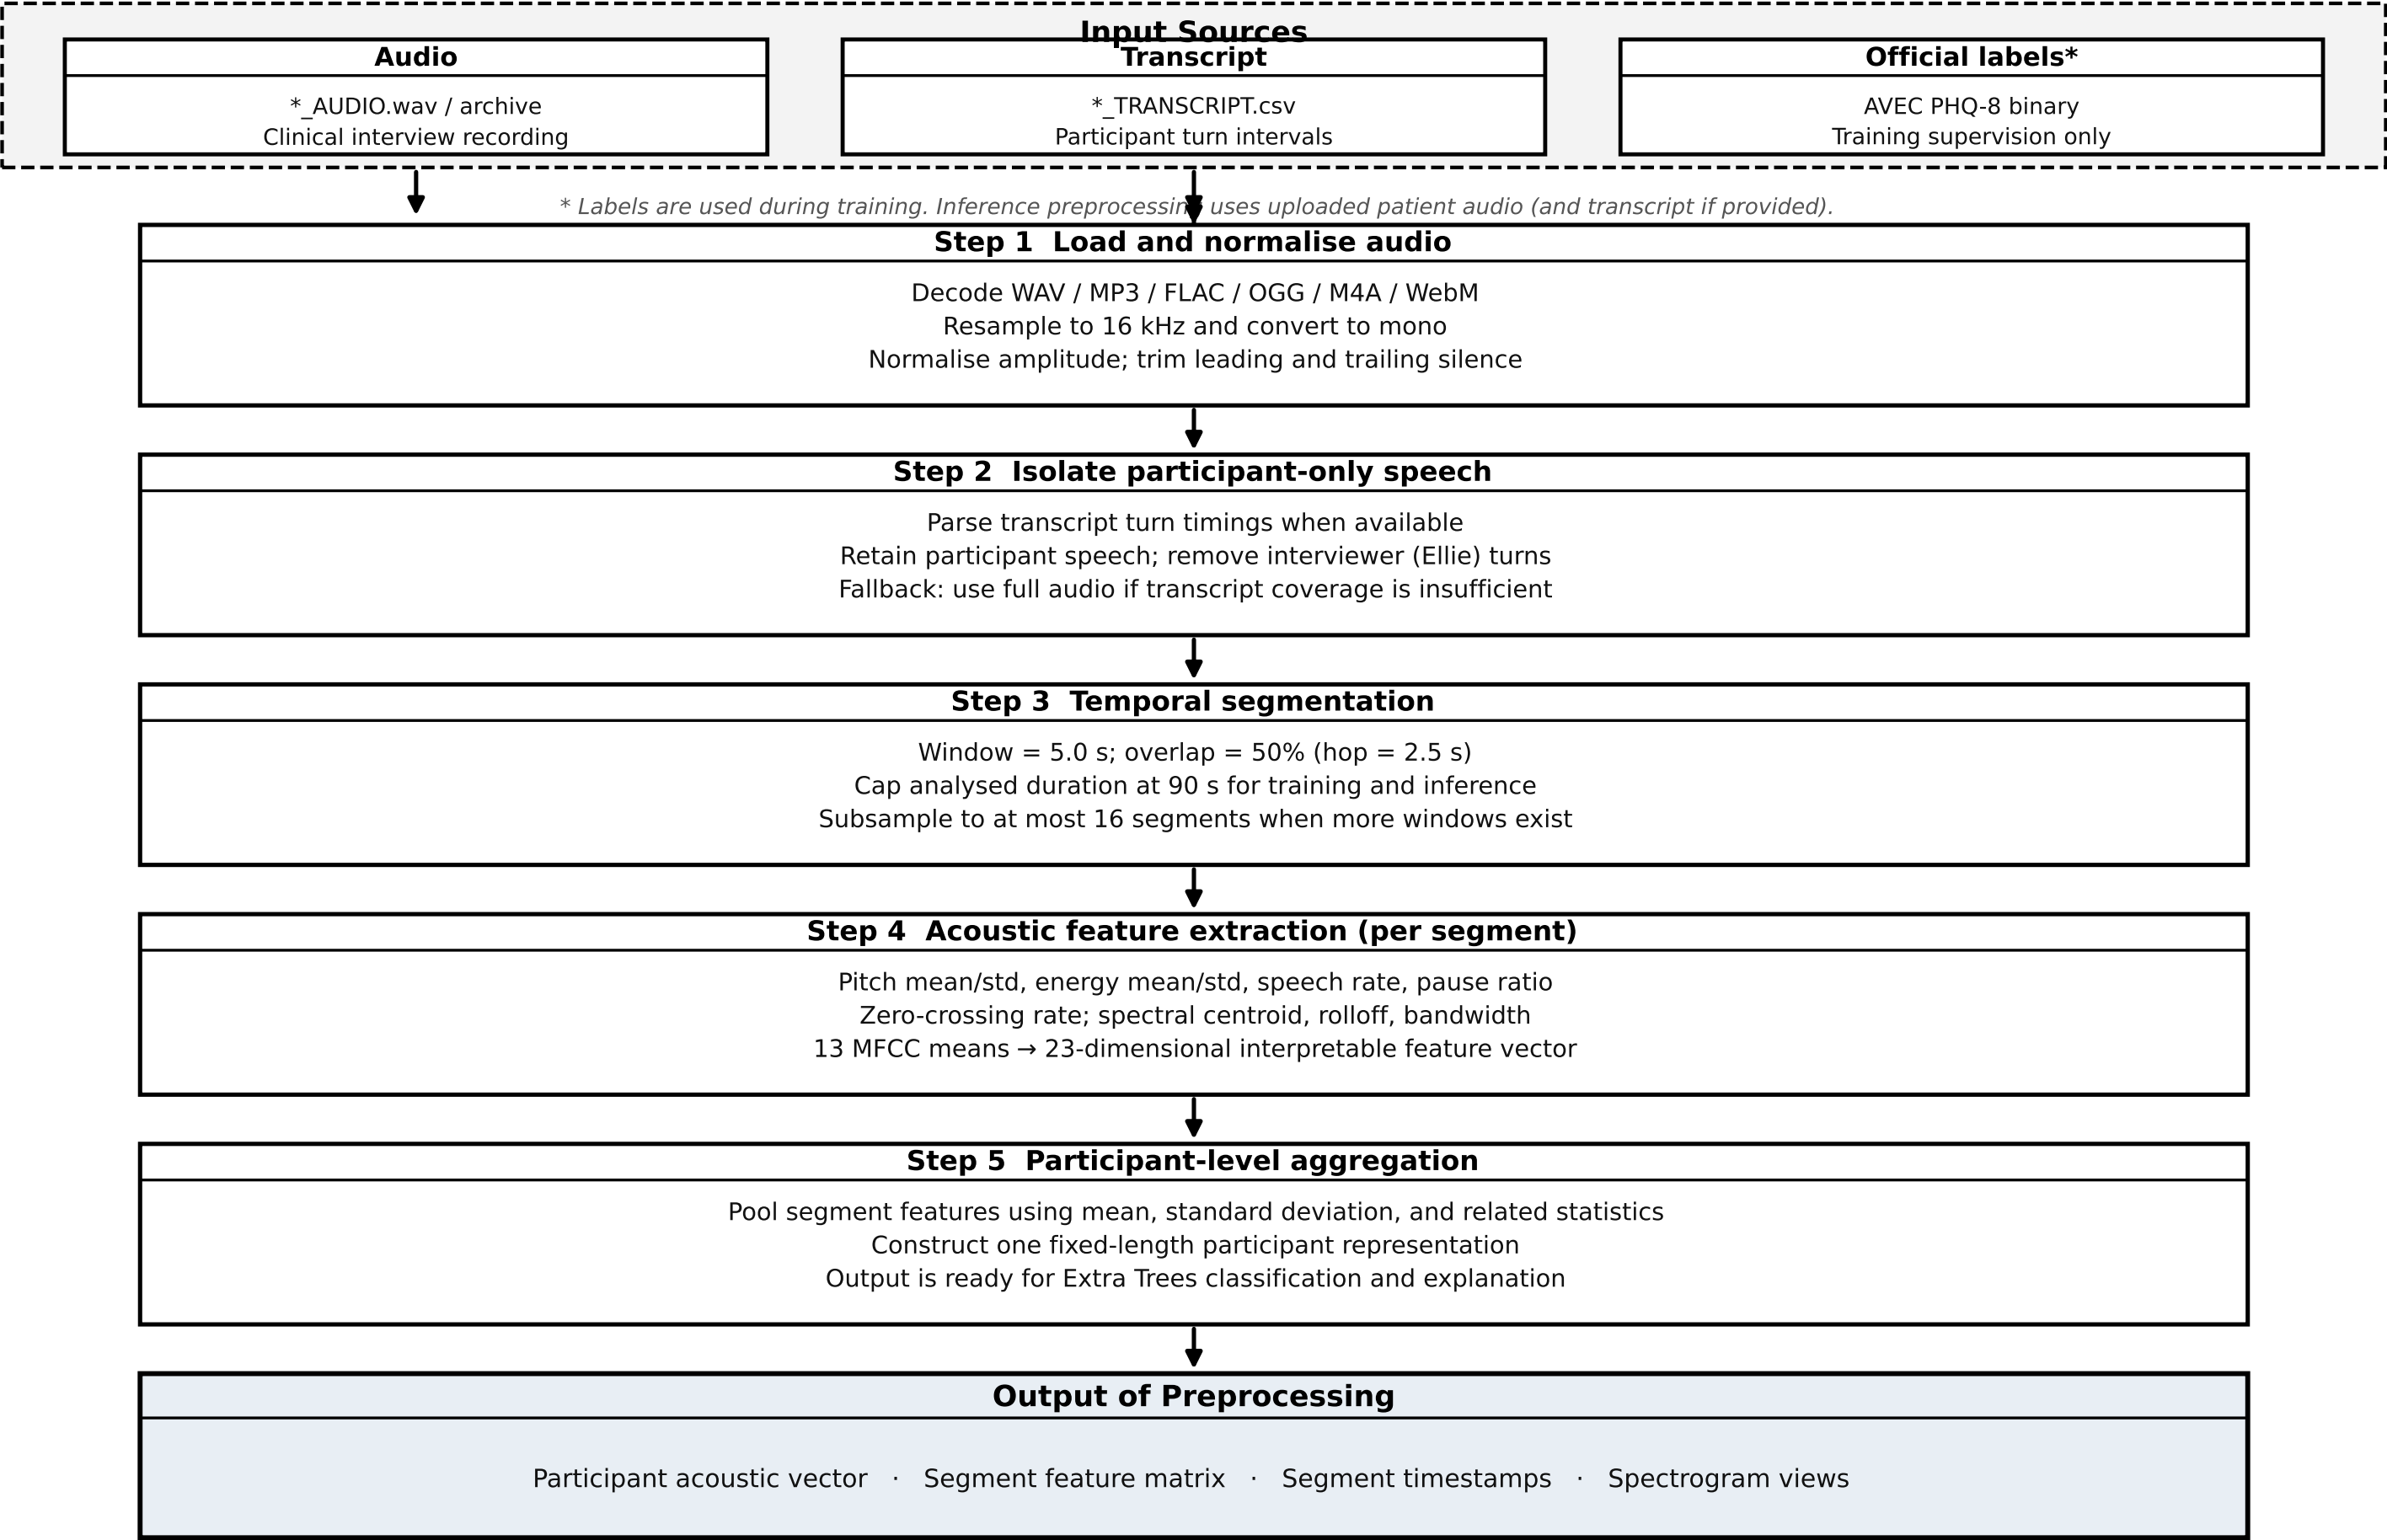


Figure 4.2 Preprocessing Pipeline

From clinical interview audio to participant-level acoustic representation



Deployed pathway (ACTIVE_MODEL = acoustic). CNN baseline uses separate 3 s / 50% mel-spectrogram segmentation.

Robustness: audio shorter than 0.5 s yields zero-filled features rather than pipeline failure.